

Machine Learning for Physics I

Artificial neural networks, optimization, and automatic differentiation

Filippo Masi[★]

Laboratoire Jean Kuntzmann (LJK), TRIPOP team

Université Grenoble Alpes, Inria, MIAI | Artificial Intelligence & Mechanics (AIM)



[★]filippo.masi@inria.fr

Outline

- 1 Learning framework and artificial neural networks
- 2 Optimizing neural networks
- 3 Automatic differentiation
- 4 Training neural networks

Learning framework and artificial neural networks

Learning framework

A learning model (e.g. a neural network) is a finite-dimensional family of functions

$$\mathbf{f}_{\boldsymbol{\theta}} : \mathcal{X} \subset \mathbb{R}^{d_X} \rightarrow \mathcal{Y} \subset \mathbb{R}^{d_Y}, \quad \text{with } \boldsymbol{\theta} \in \Theta \subset \mathbb{R}^p \text{ model parameters.}$$

Given a labeled dataset $\mathcal{D}_{XY} = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N$, with $\mathbf{x}_k \in \mathcal{X} \subset \mathbb{R}^{d_X}$ and $\mathbf{y}_k \in \mathcal{Y} \subset \mathbb{R}^{d_Y}$, the problem we aim to solve is the following problem

Learning framework

A learning model (e.g. a neural network) is a finite-dimensional family of functions

$$\mathbf{f}_{\boldsymbol{\theta}} : \mathcal{X} \subset \mathbb{R}^{d_X} \rightarrow \mathcal{Y} \subset \mathbb{R}^{d_Y}, \quad \text{with } \boldsymbol{\theta} \in \Theta \subset \mathbb{R}^p \text{ model parameters.}$$

Given a labeled dataset $\mathcal{D}_{XY} = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N$, with $\mathbf{x}_k \in \mathcal{X} \subset \mathbb{R}^{d_X}$ and $\mathbf{y}_k \in \mathcal{Y} \subset \mathbb{R}^{d_Y}$, the problem we aim to solve is the following problem

Problem: empirical loss minimization

Find $\mathbf{f}_{\boldsymbol{\theta}} : \mathcal{X} \rightarrow \mathcal{Y}$, $\mathbf{f}_{\boldsymbol{\theta}} \in \mathcal{F}_{\boldsymbol{\theta}}$ such that

$$\mathbf{f}_{\hat{\boldsymbol{\theta}}}(\mathcal{D}_{XY}) = \arg \min_{\mathbf{f}_{\boldsymbol{\theta}} \in \mathcal{F}_{\boldsymbol{\theta}}} L_{\mathcal{D}_{XY}}(\mathbf{f}_{\boldsymbol{\theta}}), \quad L_{\mathcal{D}_{XY}}(\mathbf{f}_{\boldsymbol{\theta}}) = \frac{1}{N} \sum_{i=1}^N \ell(\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}_i), \mathbf{y}_i), \quad \ell(\mathbf{f}_{\boldsymbol{\theta}}(\mathbf{x}), \mathbf{y}) \geq 0$$

where $\mathcal{F}_{\boldsymbol{\theta}}$ is the class of regression functions, $L_{\mathcal{D}_{XY}}$ the empirical loss.

Learning framework

A learning model (e.g. a neural network) is a finite-dimensional family of functions

$$\mathbf{f}_\theta : \mathcal{X} \subset \mathbb{R}^{d_X} \rightarrow \mathcal{Y} \subset \mathbb{R}^{d_Y}, \quad \text{with } \theta \in \Theta \subset \mathbb{R}^p \text{ model parameters.}$$

Given a labeled dataset $\mathcal{D}_{XY} = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N$, with $\mathbf{x}_k \in \mathcal{X} \subset \mathbb{R}^{d_X}$ and $\mathbf{y}_k \in \mathcal{Y} \subset \mathbb{R}^{d_Y}$, the problem we aim to solve is the following problem

Problem: empirical loss minimization

Find $\mathbf{f}_\theta : \mathcal{X} \rightarrow \mathcal{Y}$, $\mathbf{f}_\theta \in \mathcal{F}_\theta$ such that

$$\mathbf{f}_{\hat{\theta}}(\mathcal{D}_{XY}) = \arg \min_{\mathbf{f}_\theta \in \mathcal{F}_\theta} L_{\mathcal{D}_{XY}}(\mathbf{f}_\theta), \quad L_{\mathcal{D}_{XY}}(\mathbf{f}_\theta) = \frac{1}{N} \sum_{i=1}^N \ell(\mathbf{f}_\theta(\mathbf{x}_i), \mathbf{y}_i), \quad \ell(\mathbf{f}_\theta(\mathbf{x}), \mathbf{y}) \geq 0$$

where $\mathcal{F}_\theta$ is the class of regression functions, $L_{\mathcal{D}_{XY}}$ the empirical loss.

- ▶ $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ is the "distance" between "true" labels and predictions, e.g.
- ▶ $\ell(\mathbf{f}_\theta(\mathbf{x}_i), \mathbf{y}_i) = \|\mathbf{f}_\theta(\mathbf{x}_i) - \mathbf{y}_i\|_2^2$.
- ▶ $L_{\mathcal{D}_{XY}}$ is an approximation of the "true" loss $L = \mathbb{E}_{(\mathbf{x}, \mathbf{y}) \sim p} [\ell(\mathbf{f}_\theta(\mathbf{x}), \mathbf{y})]$, but $p(\mathbf{x}, \mathbf{y})$ is not known a priori.

Learning framework

A learning model (e.g. a neural network) is a finite-dimensional family of functions

$$\mathbf{f}_{\theta} : \mathcal{X} \subset \mathbb{R}^{d_X} \rightarrow \mathcal{Y} \subset \mathbb{R}^{d_Y}, \quad \text{with } \theta \in \Theta \subset \mathbb{R}^p \text{ model parameters.}$$


Given a labeled dataset $\mathcal{D}_{XY} = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N$, with $\mathbf{x}_k \in \mathcal{X} \subset \mathbb{R}^{d_X}$ and $\mathbf{y}_k \in \mathcal{Y} \subset \mathbb{R}^{d_Y}$, the problem we aim to solve is the following problem

Problem: empirical loss minimization

Find $\mathbf{f}_{\theta} : \mathcal{X} \rightarrow \mathcal{Y}$, $\mathbf{f}_{\theta} \in \mathcal{F}_{\theta}$ such that

$$\mathbf{f}_{\hat{\theta}}(\mathcal{D}_{XY}) = \arg \min_{\mathbf{f}_{\theta} \in \mathcal{F}_{\theta}} L_{\mathcal{D}_{XY}}(\mathbf{f}_{\theta}), \quad L_{\mathcal{D}_{XY}}(\mathbf{f}_{\theta}) = \frac{1}{N} \sum_{i=1}^N \ell(\mathbf{f}_{\theta}(\mathbf{x}_i), \mathbf{y}_i), \quad \ell(\mathbf{f}_{\theta}(\mathbf{x}), \mathbf{y}) \geq 0$$

where $\mathcal{F}_{\theta}$ is the class of regression functions, $L_{\mathcal{D}_{XY}}$ the empirical loss.



Example: least squares regression

Given observations $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^N$, a linear regression model assumes that the relationship between the dependent variable y and the regressor vector $\mathbf{x}_i \in \mathbb{R}^{d_X}$ is linear, i.e.

$$\tilde{y}_i = f_{\boldsymbol{\theta}}(\mathbf{x}) = \theta_0 + \theta_1 \tilde{x}_i^{(1)} + \dots + \theta_{d_X} \tilde{x}_i^{(d_X)} = \tilde{\mathbf{x}}_i^\top \boldsymbol{\theta}, \quad \tilde{\mathbf{y}} = \mathbf{X} \boldsymbol{\theta}$$

with $\tilde{\mathbf{x}}_i = \left[1, x_i^{(1)}, x_i^{(2)}, \dots, x_i^{(d_X)}\right]^\top$ and $\mathbf{X} = \left[\mathbf{x}_1^\top, \mathbf{x}_2^\top, \dots, \mathbf{x}_N^\top\right]^\top \in \mathbb{R}^{N \times (d_X+1)}$.

The optimum parameter vector is defined as the one that minimizes the sum of mean squared (empirical) loss:

$$\hat{\boldsymbol{\theta}} = \arg \min_{\boldsymbol{\theta}} L_{\mathcal{D}}(\boldsymbol{\theta}), \quad L_{\mathcal{D}}(\boldsymbol{\theta}) = \frac{1}{2N} \|\mathbf{X}\boldsymbol{\theta} - \mathbf{y}\|_2^2, \quad \nabla_{\boldsymbol{\theta}} L_{\mathcal{D}} = \frac{1}{N} \mathbf{X}^\top (\mathbf{X}\boldsymbol{\theta} - \mathbf{y})$$

The optimum parameter vector is the solution of

$$\mathbf{X}^\top \mathbf{X} \boldsymbol{\theta} - \mathbf{X}^\top \mathbf{y} = \mathbf{0}, \quad \text{or} \quad \hat{\boldsymbol{\theta}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}.$$

Note

In a neural network, the dependence of $f_{\boldsymbol{\theta}}$ on $\boldsymbol{\theta}$ is nonlinear, so the training problem is generally nonlinear and nonconvex.

Regression viewpoint: what is being approximated?

The dataset is viewed as samples of two random variables (X, Y) drawn from an unknown joint distribution $p(x, y)$. For any random quantity Z , $\mathbb{E}[Z]$ denotes its average under this distribution. The conditional expectation $\mathbb{E}[Y | X = x]$ is the average target value among samples whose input is fixed at x .

Population target for square loss

For scalar regression, the ideal predictor is

$$f^* \in \arg \min_f \mathbb{E}[(Y - f(X))^2], \quad f^*(x) = \mathbb{E}[Y | X = x].$$

Equivalently,

$$Y = f^*(X) + \epsilon, \quad \epsilon := Y - f^*(X), \quad \mathbb{E}[\epsilon | X] = 0.$$

Here ϵ is the residual/noise random variable, not the small numerical constant used by optimizers.

What training can and cannot reduce

The conditional variance

$$\text{Var}(Y | X = x) = \text{Var}(\epsilon | X = x)$$

is irreducible at the input x . Training a neural network f_θ can only reduce the excess error due to $f_\theta \neq f^*$.

The artificial neuron

Artificial neurons have been formulated by the Pitts&McCulloch and Rosenblatt (1950s).

A neuron performs the following nonlinear transformation on its inputs $\mathbf{x}$

$$f_{\theta}(\mathbf{x}) = \mathcal{A}(a(\mathbf{x})) = \mathcal{A}(\langle \mathbf{W}, \mathbf{x} \rangle + b) = \mathcal{A}\left(\sum_{c=1}^{d_X} W_c \cdot x_c + b\right), \quad \theta := \{\mathbf{W}, b\}$$

where a is the linear pre-activation, $\mathbf{W}$ the weights, b the bias, and $\mathcal{A}$ the nonlinear activation function.

The artificial neuron

Artificial neurons have been formulated by the Pitts&McCulloch and Rosenblatt (1950s).


A neuron performs the following nonlinear transformation on its inputs $\mathbf{x}$

$$f_{\theta}(\mathbf{x}) = \mathcal{A}(a(\mathbf{x})) = \mathcal{A}(\langle \mathbf{W}, \mathbf{x} \rangle + b) = \mathcal{A}\left(\sum_{c=1}^{d_X} W_c \cdot x_c + b\right), \quad \theta := \{\mathbf{W}, b\}$$

where a is the linear pre-activation, $\mathbf{W}$ the weights, b the bias, and $\mathcal{A}$ the nonlinear activation function.

◦ List of activation functions:

– sigmoid: $\sigma(a) = \frac{1}{1 + e^{-a}}$



The artificial neuron

Artificial neurons have been formulated by the Pitts&McCulloch and Rosenblatt (1950s).

A neuron performs the following nonlinear transformation on its inputs $\mathbf{x}$


$$f_{\theta}(\mathbf{x}) = \mathcal{A}(a(\mathbf{x})) = \mathcal{A}(\langle \mathbf{W}, \mathbf{x} \rangle + b) = \mathcal{A}\left(\sum_{c=1}^{d_X} W_c \cdot x_c + b\right), \quad \theta := \{\mathbf{W}, b\}$$

where a is the linear pre-activation, $\mathbf{W}$ the weights, b the bias, and $\mathcal{A}$ the nonlinear activation function.

◦ List of activation functions:

– sigmoid: $\sigma(a) = \frac{1}{1 + e^{-a}}$

– hyperbolic tangent: $\tanh(a) = \frac{e^{2a} - 1}{e^{2a} + 1}$



The artificial neuron

Artificial neurons have been formulated by the Pitts&McCulloch and Rosenblatt (1950s).

A neuron performs the following nonlinear transformation on its inputs $\mathbf{x}$

$$f_{\theta}(\mathbf{x}) = \mathcal{A}(a(\mathbf{x})) = \mathcal{A}(\langle \mathbf{W}, \mathbf{x} \rangle + b) = \mathcal{A}\left(\sum_{c=1}^{d_X} W_c \cdot x_c + b\right), \quad \theta := \{\mathbf{W}, b\}$$


where a is the linear pre-activation, $\mathbf{W}$ the weights, b the bias, and $\mathcal{A}$ the nonlinear activation function.

◦ List of activation functions:

– sigmoid: $\sigma(a) = \frac{1}{1 + e^{-a}}$

– hyperbolic tangent: $\tanh(a) = \frac{e^{2a} - 1}{e^{2a} + 1}$

– rectified linear unit: $\text{ReLU}(a) = \max(0, a)$



The artificial neuron

Artificial neurons have been formulated by the Pitts&McCulloch and Rosenblatt (1950s).


A neuron performs the following nonlinear transformation on its inputs $\mathbf{x}$

$$f_{\theta}(\mathbf{x}) = \mathcal{A}(a(\mathbf{x})) = \mathcal{A}(\langle \mathbf{W}, \mathbf{x} \rangle + b) = \mathcal{A}\left(\sum_{c=1}^{d_X} W_c \cdot x_c + b\right), \quad \theta := \{\mathbf{W}, b\}$$

where a is the linear pre-activation, $\mathbf{W}$ the weights, b the bias, and $\mathcal{A}$ the nonlinear activation function.

◦ List of activation functions:

- sigmoid: $\sigma(a) = \frac{1}{1 + e^{-a}}$
- hyperbolic tangent: $\tanh(a) = \frac{e^{2a} - 1}{e^{2a} + 1}$
- rectified linear unit: $\text{ReLU}(a) = \max(0, a)$
- leaky ReLU: $\text{LReLU}(a) = \begin{cases} a, & a \geq 0 \\ \alpha a, & a < 0 \end{cases}$



The artificial neuron

Artificial neurons have been formulated by the Pitts&McCulloch and Rosenblatt (1950s).

A neuron performs the following nonlinear transformation on its inputs $\mathbf{x}$

$$f_{\theta}(\mathbf{x}) = \mathcal{A}(a(\mathbf{x})) = \mathcal{A}(\langle \mathbf{W}, \mathbf{x} \rangle + b) = \mathcal{A}\left(\sum_{c=1}^{d_X} W_c \cdot x_c + b\right), \quad \theta := \{\mathbf{W}, b\}$$

where a is the linear pre-activation, $\mathbf{W}$ the weights, b the bias, and $\mathcal{A}$ the nonlinear activation function.

◦ List of activation functions:


– sigmoid: $\sigma(a) = \frac{1}{1 + e^{-a}}$

– hyperbolic tangent: $\tanh(a) = \frac{e^{2a} - 1}{e^{2a} + 1}$

– rectified linear unit: $\text{ReLU}(a) = \max(0, a)$

– leaky ReLU: $\text{LReLU}(a) = \begin{cases} a, & a \geq 0 \\ \alpha a, & a < 0 \end{cases}$







– exponential linear unit: $\text{ELU}(a) = \begin{cases} a, & a \geq 0 \\ \alpha (e^a - 1), & a < 0 \end{cases}$



Activation functions

List of activation functions (non-exhaustive):

sigmoid, hyperbolic tangent, ReLU, LReLU, ELU, Softplus ($= \ln(1 + e^a)$).



The multilayer perceptron (MLP)

Stacking two layers of neurons yields an extremely flexible predictor model (MLP)

$$f_{\theta}(\mathbf{x}) = \mathcal{A}(a(\phi(\mathbf{x}))) = \mathcal{A}(\langle \mathbf{W}, \phi(\mathbf{x}) \rangle + b) = \mathcal{A}\left(\sum_{c=1}^{d_X} W_c \cdot \phi_c(\mathbf{x}) + b\right), \quad \theta := \{\mathbf{W}, b\}$$


where $\phi : \mathbf{x} \rightarrow \mathbb{R}^{d_\phi}$ is an intermediate transformation (hidden layer) mapping $\mathbf{x}$ into the space of hidden features of d_ϕ .

The multilayer perceptron (MLP)

Stacking two layers of neurons yields an extremely flexible predictor model (MLP)

$$f_{\theta}(\mathbf{x}) = \mathcal{A}(a(\phi(\mathbf{x}))) = \mathcal{A}(\langle \mathbf{W}, \phi(\mathbf{x}) \rangle + b) = \mathcal{A}\left(\sum_{c=1}^{d_X} W_c \cdot \phi_c(\mathbf{x}) + b\right), \quad \theta := \{\mathbf{W}, b\}$$

where $\phi: \mathbf{x} \rightarrow \mathbb{R}^{d_\phi}$ is an intermediate transformation (hidden layer) mapping $\mathbf{x}$ into the space of hidden features of d_ϕ .



The neural network output reads (layers: one hidden $\mathbf{f}^{(1)}$ and output $\mathbf{f}^{(o)}$)

$$f_{\theta}(\mathbf{x}) = f^{(o)} \circ \mathbf{f}^{(1)}(\mathbf{x}) = \mathcal{A}^{(o)}(a^{(o)}(\mathbf{f}^{(1)}(\mathbf{x})))$$

$$= \mathcal{A}^{(o)}\left(\langle \mathbf{W}^{(o)}, \mathbf{f}^{(1)}(\mathbf{x}) \rangle + b^{(o)}\right) = \mathcal{A}^{(o)}\left(\sum_{c=1}^{d_{f_1}} W_c^{(o)} \cdot f_c^{(1)}(\mathbf{x}) + b^{(o)}\right)$$

$$\text{with } f_c^{(1)} = \mathcal{A}^{(1)}(a_c^{(1)}(\mathbf{x})) = \mathcal{A}^{(1)}\left(\langle \mathbf{W}^{(1)}, \mathbf{x} \rangle + b^{(1)}\right) = \mathcal{A}^{(1)}\left(\sum_{c=1}^{d_X} W_c^{(1)} \cdot x_c + b^{(1)}\right)$$

Stacking MLP: artificial neural networks

Using a MLP and stacking multiple hidden layers yields a (feed-forward) artificial neural network:

$$\begin{aligned}\mathbf{f}^{(0)} &= \mathbf{x}, \\ \mathbf{a}^{(\ell+1)} &= \langle \mathbf{W}^{(\ell)}, \mathbf{f}^{(\ell)} \rangle + \mathbf{b}^{(\ell)}, \\ \mathbf{f}^{(\ell+1)} &= \mathcal{A}_\ell \left(\mathbf{a}^{(\ell+1)} \right), && \text{hidden layer } \ell = 0, \dots, L-1, \\ \mathbf{f}_\theta(\mathbf{x}) &= \langle \mathbf{W}^{(o)}, \mathbf{f}^{(o)} \rangle + \mathbf{b}^{(o)}, && \text{output layer } (o),\end{aligned}$$

where the parameter vector is $\theta = \{ \mathbf{W}^{(\ell)}, \mathbf{b}^{(\ell)}, \mathbf{W}^{(o)}, \mathbf{b}^{(o)} \}$.

Stacking MLP: artificial neural networks

Using a MLP and stacking multiple hidden layers yields a (feed-forward) artificial neural network:

$$\begin{aligned}\mathbf{f}^{(0)} &= \mathbf{x}, \\ \mathbf{a}^{(\ell+1)} &= \langle \mathbf{W}^{(\ell)}, \mathbf{f}^{(\ell)} \rangle + \mathbf{b}^{(\ell)}, \\ \mathbf{f}^{(\ell+1)} &= \mathcal{A}_\ell \left(\mathbf{a}^{(\ell+1)} \right), & \text{hidden layer } \ell = 0, \dots, L-1, \\ \mathbf{f}_\theta(\mathbf{x}) &= \langle \mathbf{W}^{(o)}, \mathbf{f}^{(o)} \rangle + \mathbf{b}^{(o)}, & \text{output layer } (o),\end{aligned}$$

where the parameter vector is $\theta = \{ \mathbf{W}^{(\ell)}, \mathbf{b}^{(\ell)}, \mathbf{W}^{(o)}, \mathbf{b}^{(o)} \}$.


How many layers? Lower bounds for approximation by MLP neural nets

Given a function $h : [-T, T]^{d_X} \subset \mathcal{X} \rightarrow \mathbb{R}$, $T > 0$, $h \in \mathcal{W}^{k,2}$, there exists a one-hidden-layer MLP $\hat{f}_\theta$ with sigmoid activation function and N_K hidden neurons such that the L^2 approximation error $e := \|h - \hat{f}_\theta\|_{L^2([-T, T]^{d_X})}$ is smaller than a prescribed tolerance ϵ if

$$N_K \approx \epsilon^{\frac{1-d_X}{k}} = \epsilon^{-\frac{d_X-1}{k}}, \quad d_X = \dim(\mathcal{X})$$

For $h : L^2[-T, T]^{d_X}$, to reduce the error by an order or magnitude, the number of neurons (width) must be multiplied by $10^{\frac{d_X-1}{2}}$

Stacking MLP: artificial neural networks



Expressivity and derivatives

Neural networks are used in scientific machine learning because they can represent flexible functions while remaining differentiable programs.

Universal approximation, operational view

With enough hidden units, a one-hidden-layer network with a suitable nonlinearity can approximate any continuous target on a compact set to prescribed accuracy. Depth and width control how this approximation is realized in practice.

For PDE-constrained learning, the network output is often interpreted as a field

$$\mathbf{u}_\theta : \Omega \subset \mathbb{R}^d \rightarrow \mathbb{R}^q,$$

and the loss may require $\nabla \mathbf{u}_\theta$, Hessians, or higher-order derivatives. This is why the next two ingredients are central:

- ▶ optimization, to fit the parameters θ ;
- ▶ automatic differentiation, to compute derivatives of $\mathbf{u}_\theta$ and of the loss.

Optimizing neural networks

Gradient descent

Training seeks parameters $\hat{\theta} \in \arg \min_{\theta \in \Theta} L_{\mathcal{D}_{XY}}(\theta)$.

Throughout this section, the superscript (i) denotes the current optimization epoch/iteration. Starting from $\theta^{(0)}$, define

$$\mathbf{g}^{(i)} := \nabla_{\theta} L_{\mathcal{D}_{XY}}(\theta^{(i)}).$$

Gradient descent updates

$$\theta^{(i+1)} = \theta^{(i)} - \eta^{(i)} \mathbf{g}^{(i)}, \quad \eta^{(i)} > 0.$$

Interpretation

The gradient points in the direction of steepest local increase of the loss. Its negative therefore points in the direction of steepest local decrease.

For a suitable learning rate,

$$L_{\mathcal{D}_{XY}}(\theta^{(i+1)}) < L_{\mathcal{D}_{XY}}(\theta^{(i)})$$

whenever $\mathbf{g}^{(i)} \neq \mathbf{0}$, at least locally.

Why gradient descent?

Let $L_{\mathcal{D}} : \Theta \rightarrow \mathbb{R}$ be differentiable and let $\Theta \subset \mathbb{R}^p$ be convex. If $\hat{\theta}$ is a local minimizer, then the first-order necessary condition is

$$\langle \nabla_{\theta} L_{\mathcal{D}}(\hat{\theta}), \theta - \hat{\theta} \rangle \geq 0, \quad \forall \theta \in \Theta.$$

If $L_{\mathcal{D}}$ is convex, this condition is also sufficient for a global minimizer.

At a current iterate $\theta^{(i)}$, a first-order Taylor expansion gives

$$L_{\mathcal{D}}(\theta^{(i)} + \delta\theta) = L_{\mathcal{D}}(\theta^{(i)}) + \langle \nabla_{\theta} L_{\mathcal{D}}(\theta^{(i)}), \delta\theta \rangle + o(\|\delta\theta\|).$$

Among all perturbations of fixed norm $\|\delta\theta\| = r$,

$$\delta\theta^{\star} = -r \frac{\nabla_{\theta} L_{\mathcal{D}}(\theta^{(i)})}{\|\nabla_{\theta} L_{\mathcal{D}}(\theta^{(i)})\|}.$$

Steepest-descent update

Therefore, in the unconstrained case,

$$\theta^{(i+1)} = \theta^{(i)} - \eta^{(i)} \nabla_{\theta} L_{\mathcal{D}}(\theta^{(i)}).$$

If the gradient is nonzero, a sufficiently small $\eta^{(i)} > 0$ decreases the loss.

Stochastic Gradient Descent

The full empirical gradient at epoch i is

$$\mathbf{g}^{(i)} = \frac{1}{N} \sum_{n=1}^N \nabla_{\theta} \ell(\mathbf{f}_{\theta^{(i)}}(\mathbf{x}_n), \mathbf{y}_n) .$$

Stochastic Gradient Descent

The full empirical gradient at epoch i is

$$\mathbf{g}^{(i)} = \frac{1}{N} \sum_{n=1}^N \nabla_{\boldsymbol{\theta}} \ell(\mathbf{f}_{\boldsymbol{\theta}^{(i)}}(\mathbf{x}_n), \mathbf{y}_n) .$$

◦ Gradient descent

$$\boldsymbol{\theta}^{(i+1)} = \boldsymbol{\theta}^{(i)} - \eta^{(i)} \mathbf{g}^{(i)}$$


Stochastic Gradient Descent

The full empirical gradient at epoch i is

$$\mathbf{g}^{(i)} = \frac{1}{N} \sum_{n=1}^N \nabla_{\boldsymbol{\theta}} \ell(\mathbf{f}_{\boldsymbol{\theta}^{(i)}}(\mathbf{x}_n), \mathbf{y}_n) .$$

◦ Gradient descent

$$\boldsymbol{\theta}^{(i+1)} = \boldsymbol{\theta}^{(i)} - \eta \mathbf{g}^{(i)}$$




Stochastic Gradient Descent

The full empirical gradient at epoch i is

$$\mathbf{g}^{(i)} = \frac{1}{N} \sum_{n=1}^N \nabla_{\boldsymbol{\theta}} \ell(\mathbf{f}_{\boldsymbol{\theta}^{(i)}}(\mathbf{x}_n), \mathbf{y}_n) .$$

◦ Gradient descent

$$\boldsymbol{\theta}^{(i+1)} = \boldsymbol{\theta}^{(i)} - \eta \mathbf{g}^{(i)}$$



Stochastic Gradient Descent

The full empirical gradient at epoch i is



$$\mathbf{g}^{(i)} = \frac{1}{N} \sum_{n=1}^N \nabla_{\theta} \ell(\mathbf{f}_{\theta^{(i)}}(\mathbf{x}_n), \mathbf{y}_n) .$$

◦ Gradient descent

$$\theta^{(i+1)} = \theta^{(i)} - \eta^{(i)} \mathbf{g}^{(i)}$$

◦ Stochastic gradient descent

$$\theta^{(i+1)} = \theta^{(i)} - \eta^{(i)} \nabla_{\theta} \ell(\mathbf{f}_{\theta^{(i)}}(\mathbf{x}_n), \mathbf{y}_n)$$



Stochastic Gradient Descent

The full empirical gradient at epoch i is



$$\mathbf{g}^{(i)} = \frac{1}{N} \sum_{n=1}^N \nabla_{\theta} \ell(\mathbf{f}_{\theta^{(i)}}(\mathbf{x}_n), \mathbf{y}_n) .$$

◦ Gradient descent

$$\theta^{(i+1)} = \theta^{(i)} - \eta^{(i)} \mathbf{g}^{(i)}$$

◦ Stochastic gradient descent

$$\theta^{(i+1)} = \theta^{(i)} - \eta^{(i)} \nabla_{\theta} \ell(\mathbf{f}_{\theta^{(i)}}(\mathbf{x}_n), \mathbf{y}_n)$$




SGD algorithm

Given $\mathcal{D}_{XY} = \{(\mathbf{x}_n, \mathbf{y}_n)\}_{n=1}^N$, we optimize the parameters using mini-batches of size N_b .

Mini-batch stochastic gradient descent

- 1: Initialize parameters $\boldsymbol{\theta}^{(0)}$
- 2: **for** $i = 0, \dots, n_e - 1$ **do**
- 3: Randomly shuffle the indices $\{1, \dots, N\}$
- 4: Split them into mini-batches $\mathcal{B}_1, \dots, \mathcal{B}_M$, $M = \lceil N/N_b \rceil$
- 5: Set $\boldsymbol{\theta}^{(i,0)} \leftarrow \boldsymbol{\theta}^{(i)}$
- 6: **for** $m = 1, \dots, M$ **do**
- 7: Compute $\mathbf{g}^{(i,m)} \leftarrow \frac{1}{|\mathcal{B}_m|} \sum_{n \in \mathcal{B}_m} \nabla_{\boldsymbol{\theta}} \ell(\mathbf{f}_{\boldsymbol{\theta}^{(i,m-1)}}(\mathbf{x}_n), \mathbf{y}_n)$
- 8: Update $\boldsymbol{\theta}^{(i,m)} \leftarrow \boldsymbol{\theta}^{(i,m-1)} - \eta^{(i)} \mathbf{g}^{(i,m)}$
- 9: **end for**
- 10: Set $\boldsymbol{\theta}^{(i+1)} \leftarrow \boldsymbol{\theta}^{(i,M)}$
- 11: **end for**



Learning-rate scheduling

The learning rate $\eta^{(i)}$ controls the size of each parameter update. A fixed value is rarely optimal throughout training.

Too large

- ▶ unstable iterates;
- ▶ oscillation around sharp valleys;
- ▶ possible numerical overflow or NaNs.

Too small

- ▶ slow convergence;
- ▶ early apparent stagnation;
- ▶ poor practical use of compute time.

Common schedules start larger and decay later, for instance

$$\eta^{(i)} = \eta_0 0.1^{i/s} \quad \text{or} \quad \eta^{(i)} = \begin{cases} \eta_0, & i < i_1, \\ \eta_1, & i_1 \leq i < i_2, \\ \eta_2, & i \geq i_2. \end{cases}$$


Gradient descent with momentum

Let $\mathbf{g}^{(i)} = \nabla_{\theta} L_{\mathcal{D}}(\theta^{(i)})$. Momentum uses an exponentially weighted memory of previous gradients:


$$\mathbf{v}^{(i+1)} = \gamma \mathbf{v}^{(i)} + \mathbf{g}^{(i)}, \quad \theta^{(i+1)} = \theta^{(i)} - \eta^{(i)} \mathbf{v}^{(i+1)}, \quad 0 \leq \gamma < 1.$$

Equivalently, one may absorb the learning rate into the velocity update.

SGD (2000 steps, dist = 1.72)



SGD with momentum (2000 steps, dist = 0.2)



<https://sesen.ai/blog/optimisers-sgd-momentum-adam-from-scratch>

Adaptive optimizers: Adagrad and RMSProp

Let $\mathbf{g}^{(i)} = \nabla_{\boldsymbol{\theta}} L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)})$. The vector $\mathbf{s}$ stores component-wise squared-gradient history.¹

Adaptive Gradient (Adagrad)

Adapt the learning rate based on the history of its squared gradients.

Initialize $\mathbf{s}^{(0)} = \mathbf{0}$. For $i \geq 1$,

$$\mathbf{s}^{(i)} = \mathbf{s}^{(i-1)} + \mathbf{g}^{(i)} \odot \mathbf{g}^{(i)}, \quad \boldsymbol{\theta}^{(i+1)} = \boldsymbol{\theta}^{(i)} - \eta^{(i)} \frac{\mathbf{g}^{(i)}}{\sqrt{\mathbf{s}^{(i)} + \epsilon}}.$$

Root Mean Square Propagation (RMSProp)

Popular variation of Adagrad, which introduces exponential decay, so that older gradients are “forgotten”.

With $0 < \beta < 1$,

$$\mathbf{s}^{(i)} = \beta \mathbf{s}^{(i-1)} + (1 - \beta) \mathbf{g}^{(i)} \odot \mathbf{g}^{(i)}, \quad \boldsymbol{\theta}^{(i+1)} = \boldsymbol{\theta}^{(i)} - \eta^{(i)} \frac{\mathbf{g}^{(i)}}{\sqrt{\mathbf{s}^{(i)} + \epsilon}}.$$

If $\beta = 0$, RMSprop reduces to standard GD method, while $\beta \geq 1$ leads to Adagrad.

¹All products, divisions, and square roots below are component-wise. Here, $\epsilon > 0$ prevents division by zero.

Adaptive momentum optimizer: Adam

Adam is a hybrid between momentum and Adagrad and combines a momentum-like first moment with an adaptive second moment. Let $\mathbf{g}^{(i)} = \nabla_{\theta} L_{\mathcal{D}}(\theta^{(i)})$ and initialize $\mathbf{m}^{(0)} = \mathbf{s}^{(0)} = \mathbf{0}$.

$$\mathbf{m}^{(i)} = \beta_1 \mathbf{m}^{(i-1)} + (1 - \beta_1) \mathbf{g}^{(i)},$$

$$\mathbf{s}^{(i)} = \beta_2 \mathbf{s}^{(i-1)} + (1 - \beta_2) \mathbf{g}^{(i)} \odot \mathbf{g}^{(i)},$$

$$\widehat{\mathbf{m}}^{(i)} = \frac{\mathbf{m}^{(i)}}{1 - \beta_1^i}, \quad \widehat{\mathbf{s}}^{(i)} = \frac{\mathbf{s}^{(i)}}{1 - \beta_2^i},$$

$$\theta^{(i+1)} = \theta^{(i)} - \eta^{(i)} \frac{\widehat{\mathbf{m}}^{(i)}}{\sqrt{\widehat{\mathbf{s}}^{(i)} + \epsilon}}.$$

Common choices

$0 < \beta_1, \beta_2 < 1$, typically $\beta_1 = 0.9$ and $\beta_2 = 0.999$.

From gradient descent to Newton's method

At epoch i , let $\mathbf{g}^{(i)} := \nabla_{\boldsymbol{\theta}} L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)})$. Gradient descent uses only the local slope:

$$\mathbf{d}_{\text{GD}}^{(i)} = -\mathbf{g}^{(i)} \quad \implies \quad \boldsymbol{\theta}^{(i+1)} = \boldsymbol{\theta}^{(i)} + \eta^{(i)} \mathbf{d}_{\text{GD}}^{(i)}.$$

From gradient descent to Newton's method

At epoch i , let $\mathbf{g}^{(i)} := \nabla_{\boldsymbol{\theta}} L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)})$. Gradient descent uses only the local slope:

$$\mathbf{d}_{\text{GD}}^{(i)} = -\mathbf{g}^{(i)} \quad \implies \quad \boldsymbol{\theta}^{(i+1)} = \boldsymbol{\theta}^{(i)} + \eta^{(i)} \mathbf{d}_{\text{GD}}^{(i)}.$$

A second-order approximation of the loss around $\boldsymbol{\theta}^{(i)}$ is

$$L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)} + \mathbf{d}) \approx L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)}) + (\mathbf{g}^{(i)})^{\top} \mathbf{d} + \frac{1}{2} \mathbf{d}^{\top} \mathbf{H}^{(i)} \mathbf{d},$$

where $\mathbf{H}^{(i)} := \nabla_{\boldsymbol{\theta}}^2 L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)})$ is the Hessian.

From gradient descent to Newton's method

At epoch i , let $\mathbf{g}^{(i)} := \nabla_{\boldsymbol{\theta}} L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)})$. Gradient descent uses only the local slope:

$$\mathbf{d}_{\text{GD}}^{(i)} = -\mathbf{g}^{(i)} \quad \implies \quad \boldsymbol{\theta}^{(i+1)} = \boldsymbol{\theta}^{(i)} + \eta^{(i)} \mathbf{d}_{\text{GD}}^{(i)}.$$

A second-order approximation of the loss around $\boldsymbol{\theta}^{(i)}$ is

$$L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)} + \mathbf{d}) \approx L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)}) + (\mathbf{g}^{(i)})^{\top} \mathbf{d} + \frac{1}{2} \mathbf{d}^{\top} \mathbf{H}^{(i)} \mathbf{d},$$

where $\mathbf{H}^{(i)} := \nabla_{\boldsymbol{\theta}}^2 L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)})$ is the Hessian.

Minimizing this quadratic model means solving $\mathbf{H}^{(i)} \mathbf{d}_{\text{N}}^{(i)} = -\mathbf{g}^{(i)}$.

From gradient descent to Newton's method

At epoch i , let $\mathbf{g}^{(i)} := \nabla_{\boldsymbol{\theta}} L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)})$. Gradient descent uses only the local slope:

$$\mathbf{d}_{\text{GD}}^{(i)} = -\mathbf{g}^{(i)} \quad \implies \quad \boldsymbol{\theta}^{(i+1)} = \boldsymbol{\theta}^{(i)} + \eta^{(i)} \mathbf{d}_{\text{GD}}^{(i)}.$$

A second-order approximation of the loss around $\boldsymbol{\theta}^{(i)}$ is

$$L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)} + \mathbf{d}) \approx L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)}) + (\mathbf{g}^{(i)})^{\top} \mathbf{d} + \frac{1}{2} \mathbf{d}^{\top} \mathbf{H}^{(i)} \mathbf{d},$$

where $\mathbf{H}^{(i)} := \nabla_{\boldsymbol{\theta}}^2 L_{\mathcal{D}}(\boldsymbol{\theta}^{(i)})$ is the Hessian.

Minimizing this quadratic model means solving $\mathbf{H}^{(i)} \mathbf{d}_{\text{N}}^{(i)} = -\mathbf{g}^{(i)}$.

Why not use Newton's method directly?

For a network with p parameters, $\mathbf{H}^{(i)} \in \mathbb{R}^{p \times p}$. Explicitly forming and storing a dense Hessian requires $O(p^2)$ memory. Solving the Newton system can also be costly. Moreover, for a nonconvex loss, $\mathbf{H}^{(i)}$ need not be positive definite, so the Newton direction need not be a descent direction.

Broyden-Fletcher-Goldfarb-Shanno algorithm

Suppose one optimization step has produced $\boldsymbol{\theta}^{(i+1)}$ and define

$$\mathbf{s}^{(i)} := \boldsymbol{\theta}^{(i+1)} - \boldsymbol{\theta}^{(i)}, \quad \mathbf{y}^{(i)} := \mathbf{g}^{(i+1)} - \mathbf{g}^{(i)}.$$

Broyden-Fletcher-Goldfarb-Shanno algorithm

Suppose one optimization step has produced $\theta^{(i+1)}$ and define

$$\mathbf{s}^{(i)} := \theta^{(i+1)} - \theta^{(i)}, \quad \mathbf{y}^{(i)} := \mathbf{g}^{(i+1)} - \mathbf{g}^{(i)}.$$

A first-order expansion of the gradient gives $\mathbf{g}^{(i+1)} \approx \mathbf{g}^{(i)} + \mathbf{H}^{(i)} \mathbf{s}^{(i)}$, and therefore $\mathbf{y}^{(i)} \approx \mathbf{H}^{(i)} \mathbf{s}^{(i)}$. Thus each pair $(\mathbf{s}^{(i)}, \mathbf{y}^{(i)})$ contains information about the local curvature of the loss.

Broyden-Fletcher-Goldfarb-Shanno algorithm

Suppose one optimization step has produced $\theta^{(i+1)}$ and define

$$\mathbf{s}^{(i)} := \theta^{(i+1)} - \theta^{(i)}, \quad \mathbf{y}^{(i)} := \mathbf{g}^{(i+1)} - \mathbf{g}^{(i)}.$$

A first-order expansion of the gradient gives $\mathbf{g}^{(i+1)} \approx \mathbf{g}^{(i)} + \mathbf{H}^{(i)} \mathbf{s}^{(i)}$, and therefore $\mathbf{y}^{(i)} \approx \mathbf{H}^{(i)} \mathbf{s}^{(i)}$. Thus each pair $(\mathbf{s}^{(i)}, \mathbf{y}^{(i)})$ contains information about the local curvature of the loss.

Broyden-Fletcher-Goldfarb-Shanno (BFGS) algorithm

Construct an approximation $\mathbf{B}^{(i)} \approx (\mathbf{H}^{(i)})^{-1}$ such that $\mathbf{B}^{(i+1)} \mathbf{y}^{(i)} = \mathbf{s}^{(i)}$. With $\rho^{(i)} := 1/((\mathbf{y}^{(i)})^\top \mathbf{s}^{(i)})$,

$$\begin{aligned} \mathbf{B}^{(i+1)} &= \left(\mathbf{I} - \rho^{(i)} \mathbf{s}^{(i)} (\mathbf{y}^{(i)})^\top \right) \mathbf{B}^{(i)} \left(\mathbf{I} - \rho^{(i)} \mathbf{y}^{(i)} (\mathbf{s}^{(i)})^\top \right) + \rho^{(i)} \mathbf{s}^{(i)} (\mathbf{s}^{(i)})^\top, \\ \mathbf{d}^{(i+1)} &= -\mathbf{B}^{(i+1)} \mathbf{g}^{(i+1)}. \end{aligned}$$

Broyden-Fletcher-Goldfarb-Shanno algorithm

Suppose one optimization step has produced $\theta^{(i+1)}$ and define

$$\mathbf{s}^{(i)} := \theta^{(i+1)} - \theta^{(i)}, \quad \mathbf{y}^{(i)} := \mathbf{g}^{(i+1)} - \mathbf{g}^{(i)}.$$

A first-order expansion of the gradient gives $\mathbf{g}^{(i+1)} \approx \mathbf{g}^{(i)} + \mathbf{H}^{(i)} \mathbf{s}^{(i)}$, and therefore $\mathbf{y}^{(i)} \approx \mathbf{H}^{(i)} \mathbf{s}^{(i)}$. Thus each pair $(\mathbf{s}^{(i)}, \mathbf{y}^{(i)})$ contains information about the local curvature of the loss.

Broyden-Fletcher-Goldfarb-Shanno (BFGS) algorithm

Construct an approximation $\mathbf{B}^{(i)} \approx (\mathbf{H}^{(i)})^{-1}$ such that $\mathbf{B}^{(i+1)} \mathbf{y}^{(i)} = \mathbf{s}^{(i)}$. With $\rho^{(i)} := 1/((\mathbf{y}^{(i)})^\top \mathbf{s}^{(i)})$,

$$\begin{aligned} \mathbf{B}^{(i+1)} &= \left(\mathbf{I} - \rho^{(i)} \mathbf{s}^{(i)} (\mathbf{y}^{(i)})^\top \right) \mathbf{B}^{(i)} \left(\mathbf{I} - \rho^{(i)} \mathbf{y}^{(i)} (\mathbf{s}^{(i)})^\top \right) + \rho^{(i)} \mathbf{s}^{(i)} (\mathbf{s}^{(i)})^\top, \\ \mathbf{d}^{(i+1)} &= -\mathbf{B}^{(i+1)} \mathbf{g}^{(i+1)}. \end{aligned}$$


Limited memory algorithm (L-BFGS)

Full BFGS stores $\mathbf{B}^{(i)} \in \mathbb{R}^{p \times p}$. L-BFGS stores only the most recent m curvature pairs $\{(\mathbf{s}^{(j)}, \mathbf{y}^{(j)})\}_{j=i-m}^{i-1}$ and uses them to compute products $\mathbf{B}^{(i)} \mathbf{g}^{(i)}$.

$$\text{memory: } O(mp) \quad \text{instead of} \quad O(p^2), \quad m \ll p.$$

Graphical representation


BFGS: learning local curvature



The curvature pairs modify the metric in parameter space:

$$\mathbf{B}^{(i)} : \quad -\mathbf{g}^{(i)} \mapsto -\mathbf{B}^{(i)} \mathbf{g}^{(i)}.$$

L-BFGS: remember only recent curvature



Only $\{(\mathbf{s}^{(j)}, \mathbf{y}^{(j)})\}_{j=i-m}^{i-1}$ is retained. The action of $\mathbf{B}^{(i)}$ is reconstructed from these pairs without forming the full matrix.

Automatic differentiation

Differentiation in scientific computing

- ▶ Automatic Differentiation (AD) is the efficient calculation of derivatives of numerical algorithms (including computer programs), via the use of the chain rule of calculus implemented on a computer.
- ▶ Differentiation (derivatives, gradients, Jacobians, and Hessians) is used throughout scientific computation. It is required in numerical methods for, among other applications: nonlinear optimization (including training neural networks)
 - solution of nonlinear system of equations
 - solution of implicit equations
 - solution of ODEs/PDEs.
- ▶ There are three automatic methods to perform differentiation in scientific computation:
 - finite-difference approximation: inefficient approximation
 - symbolic differentiation: *expression swell*²
 - automatic differentiation (AD).

²Symbolic differentiation can easily produce exponentially large symbolic expressions.

Why automatic differentiation for neural-network models?

Finite differences perturb inputs or parameters and re-evaluate the program. AD differentiates the program itself.

For training

The loss is scalar,

$$L : \mathbb{R}^p \rightarrow \mathbb{R},$$

so reverse-mode AD computes $\nabla_{\theta} L$ efficiently even when p is large.

AD is exact up to floating-point arithmetic for the implemented program; it is not a finite-difference approximation of the mathematical derivative.

For scientific losses

The model output is a differentiable field,

$$\mathbf{u}_{\theta}(\mathbf{x}),$$

so AD can compute $D\mathbf{u}_{\theta}$, $D^2\mathbf{u}_{\theta}$, or mixed derivatives entering differential-equation residuals.

Finite-difference algorithm

Given a function (or numerical algorithm) $f : \mathbb{R}^n \rightarrow \mathbb{R}$, its partial derivatives at $\mathbf{x} = (x_1, x_2, \dots, x_n) \in \mathbb{R}^n$ can be approximated as

$$\frac{\partial f(\mathbf{x})}{\partial x_i} \approx \frac{f(\mathbf{x} + h\mathbf{e}_i) - f(\mathbf{x})}{h},$$

where $h > 0$ is a small step size and $\mathbf{e}_i \in \mathbb{R}^n$ is unit vector with 1 in the i -th position, and zero otherwise.

Assuming smoothness of f , Taylor series expansion gives

$$f(\mathbf{x} + h\mathbf{e}_i) = f(\mathbf{x}) + h \frac{\partial f(\mathbf{x})}{\partial x_i} + O(h^2).$$

Dividing by h and rearranging gives

$$\frac{\partial f(\mathbf{x})}{\partial x_i} - \frac{f(\mathbf{x} + h\mathbf{e}_i) - f(\mathbf{x})}{h} = O(h),$$

i.e., the approximation has truncation error $O(h)$.

- ▶ The error disappears as $h \rightarrow 0$ (in theory), but does not account for the round-off error (adding a large number x_i to a small number h), e.g. in a hypothetical 4-digit machine $10.22 + 0.001258 = 10.22$.

Approximation $\neq$ machine precision.


- ▶ It is not efficient in high dimensions (the number of function evaluations grows with n).

Automatic differentiation

- ▶ AD is a method that produces numerical derivatives for arbitrarily complex functions (or programs) at machine precision (cf. finite-difference approximations.)
- ▶ The necessary assumption is that the function or computation can be organized into a sequence of simpler primitives for which derivatives are available directly.
AD accumulates information in forward mode or reverse mode using the chain rule of calculus.

Automatic differentiation

- ▶ AD is a method that produces numerical derivatives for arbitrarily complex functions (or programs) at machine precision (cf. finite-difference approximations.)
- ▶ The necessary assumption is that the function or computation can be organized into a sequence of simpler primitives for which derivatives are available directly.
AD accumulates information in forward mode or reverse mode using the chain rule of calculus.



Automatic differentiation

- ▶ AD is a method that produces numerical derivatives for arbitrarily complex functions (or programs) at machine precision (cf. finite-difference approximations.)
- ▶ The necessary assumption is that the function or computation can be organized into a sequence of simpler primitives for which derivatives are available directly.
AD accumulates information in forward mode or reverse mode using the chain rule of calculus.

$$f(x_1, x_2) = \log(x_1) + x_1 x_2 - \sin(x_2)$$

$$s_{-1} = x_1,$$

$$s_0 = x_2,$$

$$s_1 = \log(s_{-1}),$$

$$s_2 = s_{-1} \cdot s_0,$$

$$s_3 = \sin(s_0),$$

$$s_4 = s_1 + s_2,$$

$$s_5 = s_4 - s_3 = f(x_1, x_2).$$

Forward-mode automatic differentiation

Suppose we want to compute the gradient of $f(x_1, x_2) = \log(x_1) + x_1 x_2 - \sin(x_2)$ at $(x_1, x_2) = (2, 5)$:

$$\nabla_{\mathbf{x}} f(x_1, x_2) = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2} \right).$$

Computing the two components separately requires two forward passes.

Forward-mode automatic differentiation

Suppose we want to compute the gradient of $f(x_1, x_2) = \log(x_1) + x_1 x_2 - \sin(x_2)$ at $(x_1, x_2) = (2, 5)$:

$$\nabla_{\mathbf{x}} f(x_1, x_2) = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2} \right).$$

Computing the two components separately requires two forward passes.

To compute $\partial f / \partial x_1$, seed the input tangents as $\dot{x}_1 = 1$, $\dot{x}_2 = 0$, producing the following computation (on a 4-digit machine):

forward primal trace				primal tangent trace			
$s_{-1} = x_1$	=	2		$\dot{s}_{-1} = \dot{x}_1$	=	1	
$s_0 = x_2$	=	5		$\dot{s}_0 = \dot{x}_2$	=	0	
$s_1 = \log s_{-1}$	=	0.6931		$\dot{s}_1 = \dot{s}_{-1} / s_{-1}$	=	0.5	
$s_2 = s_{-1} s_0$	=	10		$\dot{s}_2 = s_0 \dot{s}_{-1} + s_{-1} \dot{s}_0$	=	5	
$s_3 = \sin s_0$	=	-0.9589		$\dot{s}_3 = \cos(s_0) \dot{s}_0$	=	0	
$s_4 = s_1 + s_2$	=	10.69		$\dot{s}_4 = \dot{s}_1 + \dot{s}_2$	=	5.5	
$s_5 = s_4 - s_3$	=	11.65		$\dot{s}_5 = \dot{s}_4 - \dot{s}_3$	=	5.5	
$f = s_5$	=	11.65		$\partial f / \partial x_1 = \dot{s}_5$	=	5.5	

Forward-mode automatic differentiation

Suppose we want to compute the gradient of $f(x_1, x_2) = \log(x_1) + x_1 x_2 - \sin(x_2)$ at $(x_1, x_2) = (2, 5)$:

$$\nabla_{\mathbf{x}} f(x_1, x_2) = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2} \right).$$

Computing the two components separately requires two forward passes.

To compute $\partial f / \partial x_1$, seed the input tangents as $\dot{x}_1 = 1$, $\dot{x}_2 = 0$, producing the following computation (on a 4-digit machine):

forward primal trace				primal tangent trace			
$s_{-1} = x_1$	=	2		$\dot{s}_{-1} = \dot{x}_1$	=	1	
$s_0 = x_2$	=	5		$\dot{s}_0 = \dot{x}_2$	=	0	
$s_1 = \log s_{-1}$	=	0.6931		$\dot{s}_1 = \dot{s}_{-1} / s_{-1}$	=	0.5	
$s_2 = s_{-1} s_0$	=	10		$\dot{s}_2 = s_0 \dot{s}_{-1} + s_{-1} \dot{s}_0$	=	5	
$s_3 = \sin s_0$	=	-0.9589		$\dot{s}_3 = \cos(s_0) \dot{s}_0$	=	0	
$s_4 = s_1 + s_2$	=	10.69		$\dot{s}_4 = \dot{s}_1 + \dot{s}_2$	=	5.5	
$s_5 = s_4 - s_3$	=	11.65		$\dot{s}_5 = \dot{s}_4 - \dot{s}_3$	=	5.5	
$f = s_5$	=	11.65		$\partial f / \partial x_1 = \dot{s}_5$	=	5.5	

Analytically, $\partial f / \partial x_1 = 1/x_1 + x_2$, and $\partial f / \partial x_1(2, 5) = \frac{1}{2} + 5 = 5.5$ (AD is exact, in this case)

Forward-mode automatic differentiation

Suppose we want to compute the gradient of $f(x_1, x_2) = \log(x_1) + x_1 x_2 - \sin(x_2)$ at $(x_1, x_2) = (2, 5)$:

$$\nabla_{\mathbf{x}} f(x_1, x_2) = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2} \right).$$

Computing the two components separately requires two forward passes.

To compute $\partial f / \partial x_2$, seed the input tangents as $\dot{x}_1 = 0$, $\dot{x}_2 = 1$, producing the following computation (on a 4-digit machine):

forward primal trace				primal tangent trace			
$s_{-1} = x_1$	=	2		$\dot{s}_{-1} = \dot{x}_1$	=	0	
$s_0 = x_2$	=	5		$\dot{s}_0 = \dot{x}_2$	=	1	
$s_1 = \log s_{-1}$	=	0.6931		$\dot{s}_1 = \dot{s}_{-1} / s_{-1}$	=	0	
$s_2 = s_{-1} s_0$	=	10		$\dot{s}_2 = s_0 \dot{s}_{-1} + s_{-1} \dot{s}_0$	=	2	
$s_3 = \sin s_0$	=	-0.9589		$\dot{s}_3 = \cos(s_0) \dot{s}_0$	=	0.2837	
$s_4 = s_1 + s_2$	=	10.69		$\dot{s}_4 = \dot{s}_1 + \dot{s}_2$	=	2	
$s_5 = s_4 - s_3$	=	11.65		$\dot{s}_5 = \dot{s}_4 - \dot{s}_3$	=	1.716	
$f = s_5$	=	11.65		$\partial f / \partial x_1 = \dot{s}_5$	=	1.716	

Analytically, $\partial f / \partial x_2 = x_1 - \cos(x_2)$, $\partial f / \partial x_2(2, 5) = 2 - \cos(2) = \underline{1.7163378} \dots$ (machine precision)

Forward-mode automatic differentiation

Given a function $y = f(\mathbf{x})$ with $\mathbf{x} \in \mathbb{R}^n$, the computation of the gradient

$$\nabla_{\mathbf{x}} f(\mathbf{x}) = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right)$$

requires to run n separate passes of forward AD. So it looks like this has the same problem of finite-difference approximations: it scales poorly with the dimension of the input.

However, for a vector function $\mathbf{y} = \mathbf{f}(\mathbf{x})$, with $\mathbf{y} \in \mathbb{R}^m$, one is able to compute the derivatives $\nabla_{x_i} \mathbf{f}$ in a single pass.

Forward-mode automatic differentiation


Given a function $y = f(\mathbf{x})$ with $\mathbf{x} \in \mathbb{R}^n$, the computation of the gradient

$$\nabla_{\mathbf{x}} f(\mathbf{x}) = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right)$$

requires to run n separate passes of forward AD. So it looks like this has the same problem of finite-difference approximations: it scales poorly with the dimension of the input.

However, for a vector function $\mathbf{y} = \mathbf{f}(\mathbf{x})$, with $\mathbf{y} \in \mathbb{R}^m$, one is able to compute the derivatives $\nabla_{x_i} \mathbf{f}$ in a single pass.

Jacobian-vector product (JVP): input-seeded directional derivative



$$\text{JVP}_f(\mathbf{x}; \mathbf{v}) = \mathbb{J}_f(\mathbf{x}) \mathbf{v} = \left. \frac{d}{d\epsilon} \mathbf{f}(\mathbf{x} + \epsilon \mathbf{v}) \right|_{\epsilon=0}$$

Columns of the Jacobian: seeds $\mathbf{v} = \mathbf{e}_j$ return column j (full Jacobian costs $\approx n$ sweeps).

- the Jacobian $\mathbb{J}_f$ of $\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ requires n passes. If $m \gg n$ we are okay, but what if $n \gg m$?

Reverse-mode automatic differentiation

When there are many more inputs than outputs ($n > m$), the solution is to run the process backwards, from the output to the input.

Reverse-mode automatic differentiation uses the notion of adjoint

$$\bar{s}_i = \frac{\partial y}{\partial s_i}$$

to each intermediate variable s_i , where y is the function output to be differentiated. These adjoints can be computed in reverse order using the chain rule, until the adjoints $\bar{s}_{-n+1}, \dots, \bar{s}_0$ for the inputs are reached (these are the sought derivatives).

Reverse-mode automatic differentiation

When there are many more inputs than outputs ($n > m$), the solution is to run the process backwards, from the output to the input.

Reverse-mode automatic differentiation uses the notion of adjoint

$$\bar{s}_i = \frac{\partial y}{\partial s_i}$$

to each intermediate variable s_i , where y is the function output to be differentiated. These adjoints can be computed in reverse order using the chain rule, until the adjoints $\bar{s}_{-n+1}, \dots, \bar{s}_0$ for the inputs are reached (these are the sought derivatives).

Therefore, the gradient of a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ can be computed in a single backward pass, but a forward pass is still necessary to populate the values of the intermediate variables for use in the reverse pass.

Reverse-mode automatic differentiation

When there are many more inputs than outputs ($n > m$), the solution is to run the process backwards, from the output to the input.


Reverse-mode automatic differentiation uses the notion of adjoint

$$\bar{s}_i = \frac{\partial y}{\partial s_i}$$

to each intermediate variable s_i , where y is the function output to be differentiated. These adjoints can be computed in reverse order using the chain rule, until the adjoints $\bar{s}_{-n+1}, \dots, \bar{s}_0$ for the inputs are reached (these are the sought derivatives).

Therefore, the gradient of a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ can be computed in a single backward pass, but a forward pass is still necessary to populate the values of the intermediate variables for use in the reverse pass.

vector-Jacobian product (VJP): output-seeded sensitivity



$$\text{VJP}_f(\mathbf{x}; \mathbf{c}) = \mathbf{c}^\top \mathbb{J}_f(\mathbf{x}) \iff \mathbb{J}_f(\mathbf{x})^\top \mathbf{c} \text{ as a column}$$

Rows of the Jacobian: seeds $\mathbf{c} = \mathbf{e}_i$ select output i (full Jacobian costs $\approx m$ sweeps).

Reverse-mode automatic differentiation

Consider again

$$y = f(x_1, x_2) = \log(x_1) + x_1 x_2 - \sin(x_2), \quad (x_1, x_2) = (2, 5).$$

The full gradient can be computed with a single backward pass. Define the adjoint

$$\bar{s}_j := \frac{\partial y}{\partial s_j}.$$

forward primal trace			reverse adjoint trace		
$y = s_5$	=	11.65			
$s_5 = s_4 - s_3$	=	11.65	$\bar{s}_5 = \partial y / \partial s_5$	=	1
$s_4 = s_1 + s_2$	=	10.69	$\bar{s}_4 = \bar{s}_5 \partial s_5 / \partial s_4 = \bar{s}_5 \cdot 1$	=	1
$s_3 = \sin(s_0)$	=	-0.9589	$\bar{s}_3 = \bar{s}_5 \partial s_5 / \partial s_3 = \bar{s}_5 \cdot (-1)$	=	-1
$s_2 = s_{-1} s_0$	=	10	$\bar{s}_2 = \bar{s}_4 \partial s_4 / \partial s_2 = \bar{s}_4 \cdot 1$	=	1
$s_1 = \log(s_{-1})$	=	0.6931	$\bar{s}_1 = \bar{s}_4 \partial s_4 / \partial s_1 = \bar{s}_4 \cdot 1$	=	1
$s_0 = x_2$	=	5	$\bar{s}_0 = \bar{s}_2 \partial s_2 / \partial s_0 + \bar{s}_3 \partial s_3 / \partial s_0 = \bar{s}_2 s_{-1} + \bar{s}_3 \cos(s_0)$	=	1.716
$s_{-1} = x_1$	=	2	$\bar{s}_{-1} = \bar{s}_2 \partial s_2 / \partial s_{-1} + \bar{s}_1 \partial s_1 / \partial s_{-1} = \bar{s}_2 s_0 + \bar{s}_1 / s_{-1}$	=	5.5
			$\bar{x}_2 = \bar{s}_0$	=	1.7163
			$\bar{x}_1 = \bar{s}_{-1}$	=	5.5

Since $s_{-1} = x_1$ and $s_0 = x_2$,

$$\nabla f(2, 5) = (\bar{s}_{-1}, \bar{s}_0) = (5.5, 1.716).$$

JVP and VJP are different Jacobian contractions

$$\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m, \quad \nabla_{\mathbf{x}} \mathbf{f} = \mathbf{J}_{\mathbf{f}}(\mathbf{x}) \in \mathbb{R}^{m \times n}$$

JVP and VJP are different Jacobian contractions

$$\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m, \quad \nabla_{\mathbf{x}} \mathbf{f} = \mathbf{J}_{\mathbf{f}}(\mathbf{x}) \in \mathbb{R}^{m \times n}$$

$$\mathbf{J}_{\mathbf{f}} \mathbf{v}$$

JVP — forward mode

seed $\mathbf{v} \in \mathbb{R}^n$ at the input, result in $\mathbb{R}^m$
directional derivative / Jacobian columns
full Jacobian: about n sweeps

JVP and VJP are different Jacobian contractions


$$\mathbf{f} : \mathbb{R}^n \rightarrow \mathbb{R}^m, \quad \nabla_{\mathbf{x}} \mathbf{f} = \mathbf{J}_{\mathbf{f}}(\mathbf{x}) \in \mathbb{R}^{m \times n}$$

$\mathbf{J}_{\mathbf{f}} \mathbf{v}$
JVP — forward mode

seed $\mathbf{v} \in \mathbb{R}^n$ at the input, result in $\mathbb{R}^m$
directional derivative / Jacobian columns
full Jacobian: about n sweeps

$\mathbf{J}_{\mathbf{f}}^T \mathbf{c}$
VJP — reverse mode

seed $\mathbf{c} \in \mathbb{R}^m$ at the output, result in $\mathbb{R}^n$
pullback / Jacobian rows
full Jacobian: about m sweeps




Training neural networks

Backpropagation of the loss: reverse-mode AD


$$L_{\mathcal{D}_{XY}}(\theta) = \frac{1}{N} \sum_{k=1}^N \ell(\mathbf{f}_{\theta}(\mathbf{x}_k), \mathbf{y}_k), \quad \nabla_{\theta} L_{\mathcal{D}_{XY}}(\theta) = \frac{1}{N} \sum_{k=1}^N \nabla_{\theta} \ell(\mathbf{f}_{\theta}(\mathbf{x}_k), \mathbf{y}_k).$$

For a single sample, let $\tilde{\mathbf{y}}_k = \mathbf{f}_{\theta}(\mathbf{x}_k)$. By the chain rule,

$$\nabla_{\theta} \ell(\tilde{\mathbf{y}}_k, \mathbf{y}_k) = [\mathbb{J}_{\theta} \mathbf{f}_{\theta}(\mathbf{x}_k)]^{\top} \nabla_{\tilde{\mathbf{y}}} \ell(\tilde{\mathbf{y}}_k, \mathbf{y}_k), \quad \mathbb{J}_{\theta} \mathbf{f}_{\theta}(\mathbf{x}_k) := \frac{\partial \mathbf{f}_{\theta}(\mathbf{x}_k)}{\partial \theta} \in \mathbb{R}^{d_Y \times p}.$$



PyTorch setup and reproducibility



```
import torch
from torch import nn
from torch.func import jacrev, jacfwd, vmap

torch.set_default_dtype(torch.float64)
torch.manual_seed(7)


device = torch.device(
    "cuda" if torch.cuda.is_available() else "cpu"
)
```

- ▶ The random seed makes the example repeatable.
- ▶ Double precision is useful when derivatives enter a differential-equation residual.
- ▶ `torch.func` supplies composable derivative transforms.

A synthetic regression dataset

We learn a noisy sample of the smooth function

$$y(x) = \sin(\pi x) + \epsilon, \quad x \in [-1, 1], \quad \epsilon \sim \mathcal{N}(0, 0.05^2).$$



```
N = 80
x_train = 2.0 * torch.rand(N, 1, device=device) - 1.0

def target(x):
    return torch.sin(torch.pi * x)


y_clean = target(x_train)
y_train = y_clean + 0.05 * torch.randn_like(y_clean)

x_test = torch.linspace(-1.0, 1.0, 400, device=device).unsqueeze(1)
y_test = target(x_test)
```

The test grid is noise-free and contains points not used by the optimizer.

A neural network in PyTorch

```
model = nn.Sequential(  
    nn.Linear(1, 32),  
    nn.Tanh(),  
    nn.Linear(32, 32),  
    nn.Tanh(),  
    nn.Linear(32, 1),  
) .to(device)  
  
n_parameters = sum(  
    p.numel() for p in model.parameters()  
)  
print(n_parameters)                # 1153
```



Loss function and optimizer

The data loss measures the mean squared prediction error,

$$L_{\mathcal{D}_{XY}}(\theta) = \frac{1}{N} \sum_{i=1}^N |f_{\theta}(x_i) - y_i|^2.$$

```
loss_fn = nn.MSELoss()
optimizer = torch.optim.SGD(
    model.parameters(), lr=2.0e-2 # SGD, learning rate = 0.02
)


prediction = model(x_train)
loss = loss_fn(prediction, y_train)
```

- ▶ `model.parameters()` exposes the tensors optimized by SGD.
- ▶ the scalar `loss` is the output seed for reverse-mode differentiation.
- ▶ the loss graph connects every parameter to one scalar objective, which is precisely the regime where a VJP is efficient.



The training loop

```
model.train()
history = []

for step in range(4_000):
    optimizer.zero_grad(set_to_none=True)
    prediction = model(x_train)      # forward pass
    loss = loss_fn(prediction, y_train)
    loss.backward()                  # reverse-mode AD
    optimizer.step()                 # parameter update
    history.append(loss.item())
```



The call to `backward()` accumulates $\nabla_{\theta} L_{\mathcal{D}_{XY}}$ in the `.grad` field of each parameter. We need to clear those fields before the next iteration to prevent unwanted accumulation.



Inspect the trained neural network

```
model.eval()

with torch.no_grad():
    y_test_pred = model(x_test)
    rmse = torch.sqrt(
        loss_fn(y_test_pred, y_test)
    )

assert y_test_pred.shape == y_test.shape
assert torch.isfinite(rmse)
print(f"test RMSE: {rmse.item():.3e}")
```

Evaluation mode


`model.eval()` fixes the behavior of layers such as dropout and batch normalization. It does not disable differentiation by itself.

No derivative graph

`torch.no_grad()` avoids recording operations during this prediction-only check.

Before differentiating the network

Leave the `no_grad` context. Input derivatives require PyTorch to record the operations from x to $f_{\theta}(x)$.



Gradients of a neural network

```
x_phys = torch.linspace(-1.0, 1.0, 200,  
                        device=device).unsqueeze(1)  
x_phys.requires_grad_(True)  
  
u = model(x_phys)  
u_x = torch.autograd.grad(  
    u.sum(), x_phys,  
    create_graph=True  
)[0]
```


```
def u_single(x):  
    return model(x).squeeze()  
  
u_x_func = vmap(jacrev(u_single))(x_phys.detach())  
  
torch.testing.assert_close(  
    u_x.detach(), u_x_func  
)
```

autograd.grad: the scalar `u.sum()` supplies a cotangent of ones.

jacrev and **vmap**: `jacrev` differentiates one scalar network evaluation, and `vmap` applies that derivative function to all coordinate points.

First-order gradient

$$\nabla_x f_\theta(x) \stackrel{?}{=} \pi \cos(\pi x)$$




Higher-order gradients of a neural network

The derivative is itself a differentiable program and differentiating it again gives the Hessian.

Second-order gradient

$$\nabla_{xx} f_{\theta}(x) \stackrel{?}{=} -\pi^2 \sin(\pi x)$$



```
grad_u = jacrev(u_single)
hess_u = jacfwd(grad_u)                                # forward over reverse

X = x_phys.detach()
u_value = vmap(u_single)(X)                             # (N,)
H = vmap(hess_u)(X)                                     # (N, 1, 1)
u_xx = H[:, 0, 0]
```

References and further reading

- ▶ A. G. Baydin, B. A. Pearlmutter, A. A. Radul, and J. M. Siskind, “Automatic Differentiation in Machine Learning: a Survey,” *Journal of Machine Learning Research*, 18:1–43, 2018.
- ▶ L. B. Rall, “Perspectives on Automatic Differentiation: Past, Present, and Future?” in *Automatic Differentiation: Applications, Theory, and Implementations*, Springer, 2006.
- ▶ F. Gatti. (2024). Artificial Neural Networks: Learning the Optimum Statistical Model from Data. In *Machine Learning in Geomechanics 1* (eds I. Stefanou and F. Darve). url: <https://doi.org/10.1002/9781394325665.ch5>.
- ▶ F. Masi (2024). Introduction to Regression Methods. In *Machine Learning in Geomechanics 1* (eds I. Stefanou and F. Darve). url: <https://doi.org/10.1002/9781394325634.ch2>.
- ▶ F. Gatti. (2024). Artificial Neural Networks: Basic architectures and training strategies. In *Machine Learning in Geomechanics 2* (eds I. Stefanou and F. Darve). url: <https://doi.org/10.1002/9781394325665.ch5>.
- ▶ I. Stefanou, F. Darve (2023). ALERT Doctoral School on Machine learning for Geomechanics. url: <https://github.com/alert-geomaterials/2023-doctoral-school..>
- ▶ U. Braga-Neto, *ECEN 744: Scientific Machine Learning, Lecture 3: Automatic Differentiation*, Spring 2026. url: sciml.tamids.tamu.edu/ecen-744-scientific-machine-learning.
- ▶ U. Braga-Neto, *ECEN 744: Scientific Machine Learning, Lecture 4: Physics-Informed Neural Networks*, Spring 2026. url: sciml.tamids.tamu.edu/ecen-744-scientific-machine-learning.